

Assignment 4 Generative Adversarial Nets (Unconditional, 10 pts)

In this exercise, we will implement a Generative Adversarial Net (GAN), specifically, a Wasserstein GAN and train it on the MNIST dataset. *We recommend completing this assignment using Google Colab with Chrome Browser.*

Submit

- 1. (Doc A) Include the two figures at the end in the pdf generated by the latex file with Exercise 2
- 2. (Doc B) The completed *.ipynb file with all the command outputs (can be created by saving the file after finishing the experiment and downloading it from Colab)

Setup

- 1. In Colab, open tab Runtime > Change runtime type, choose *python3* and *T4 GPU*.
- 2. Run the following command to set up the environment. (Takes ~ 1.5 min)

```
! pip install --quiet "ipython[notebook]==7.34.0, <8.17.0" "setuptools>=68.0.0, <68.3.0" "torch==1.13.0" "matplotlib" "torchvision"
```

Let's start with importing our standard set of libraries.

```
import time
import torch
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import torchvision.utils as vutils
from torch import nn, optim, autograd
from torch.utils.data import DataLoader
from dataclasses import dataclass

%matplotlib inline
torch.set_num_threads(1)
torch.manual_seed(1)

device = torch.device("cuda:0") if torch.cuda.is_available() else torch.device("cpu")

if device == torch.device("cuda:0"):
    print('Everything looks good; continue')
else:
    # It is OK if you cannot connect to a GPU. In this case, training the model for
    # 2 epoch is sufficient to get full mark. (NOTE THAT 2 epoch takes approximately 1.5 hours to train for CPU)
    print('GPU is not detected. Make sure you have chosen the right runtime type')

🔄 Everything looks good; continue
```

Dataloaders and hyperparameters (0 pt)

```
@dataclass
class Hyperparameter:
    batchsize: int = 64
    num_epochs: int = 5
    latent_size: int = 32
    n_critic: int = 5
    critic_size: int = 1024
    generator_size: int = 1024
    critic_hidden_size: int = 1024
    gp_lambda: float = 10.

hp = Hyperparameter()

transform = transforms.Compose(
    [transforms.ToTensor(), transforms.Normalize((0.5,), (0.5,))])

dataset = torchvision.datasets.MNIST(
    "mnist", download=True, transform=transform)
dataloader = DataLoader(dataset, batch_size=hp.batchsize,
                        num_workers=1, shuffle=True, drop_last=True, pin_memory=True)
```

Building Models (2 pts)

After examining the preprocessing steps, we can now start building the models, including the generator for generating new images from random noise, and a critic of the realness of the image.

In this assignment we adopt the implementation of [DCGAN](#), which is a direct extension of [GAN](#), with convolutional and convolutional-transpose layers in the critic and genrator, respectively. Specifically, we will use the [ConvTranspose2d](#) layers to upscale the noise.

Moreover, we apply an improved version of [Wasserstein-GAN](#) with a [Gradient Penalty](#) (you may read Algorithm 1 to fully understand the code we are implementing).

```
# Define the generator
class Generator(nn.Module):
    def __init__(self):
        super(Generator, self).__init__()

        # Add latent embedding layer to adjust the dimension of the input (1 pt)
        self.latent_embedding = nn.Linear(hp.latent_size, hp.generator_size)

        # Transposed CNN layers to transfer noise to image
        self.tcnn = nn.Sequential(
            # input is Z, going into a convolution
            nn.ConvTranspose2d(hp.generator_size, hp.generator_size, 4, 1, 0),
            nn.BatchNorm2d(hp.generator_size),
            nn.ReLU(inplace=True),
            # upscaling
            nn.ConvTranspose2d(hp.generator_size,
                               hp.generator_size // 2, 3, 2, 1),
            nn.BatchNorm2d(hp.generator_size // 2),
            nn.ReLU(inplace=True),
            # upscaling
            nn.ConvTranspose2d(hp.generator_size // 2,
                               hp.generator_size // 4, 4, 2, 1),
            nn.BatchNorm2d(hp.generator_size // 4),
            nn.ReLU(inplace=True),
            nn.ConvTranspose2d(hp.generator_size // 4, 1, 4, 2, 1),
            nn.Tanh()
        )

    def forward(self, latent):
        vec_latent = self.latent_embedding(latent).reshape(-1, hp.generator_size, 1, 1)
        return self.tcnn(vec_latent)

# Define the critic
class Critic(nn.Module):
    def __init__(self):
        super(Critic, self).__init__()

        # CNN layers that perform downscaling
        self.cnn_net = nn.Sequential(
            nn.Conv2d(1, hp.critic_size // 4, 3, 2),
            nn.InstanceNorm2d(hp.critic_size // 4, affine=True),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(hp.critic_size // 4, hp.critic_size // 2, 3, 2),
            nn.InstanceNorm2d(hp.critic_size // 2, affine=True),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(hp.critic_size // 2, hp.critic_size, 3, 2),
            nn.InstanceNorm2d(hp.critic_size, affine=True),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Flatten(),
        )

        # Linear layers that produce the output from the features
        self.critic_net = nn.Sequential(
            nn.Linear(hp.critic_size * 4, hp.critic_hidden_size),
            nn.LeakyReLU(0.2, inplace=True),

            # Add the last layer to reflect the output (1 pt)
            nn.Linear(hp.critic_hidden_size, 1)
        )

    def forward(self, image):
        cnn_features = self.cnn_net(image)
        return self.critic_net(cnn_features)
```

Before Training

Next we define the two models and the optimizers. We use the [AdamW](#) algorithm.

```
critic, generator = Critic().to(device), Generator().to(device)

critic_optimizer = optim.AdamW(critic.parameters(), lr=1e-4,betas=(0., 0.9))
generator_optimizer = optim.AdamW(generator.parameters(), lr=1e-4,betas=(0., 0.9))
```

Training pipeline (6 points)

Finally, we perform training on the two networks. The training consists of two steps: (1) Updating discriminators for `n_critic` steps (such that we have an optimal critic): here we use an aggregation of three loss functions, (a) The real loss (the output scalar of the critic for real images); (b) The fake loss (same value for fake images); (c) The [gradient penalty](#). (2) Updating generators by only considering the fake loss (to fool the critic).

```
img_list, generator_losses, critic_losses = [], [], []
iters = 0
fixed_noise = torch.randn((64, hp.latent_size), device=device)
grad_tensor = torch.ones((hp.batchsize, 1), device=device)
start_time = time.time()
for epoch in range(hp.num_epochs):
    for batch_idx, data in enumerate(dataloader, 0):
        real_images = data[0].to(device)

        # Update Critic
        critic_optimizer.zero_grad()

        # (a) Real loss
```

```

    # (a) Real loss
    critic_output_real = critic(real_images)
    critic_loss_real = critic_output_real.mean()

# (b) Fake loss
# (1) Generating a noise tensor (of dimension (batch_size, latent_size)), you are required to
# use the hyperparameters in the hp class (0.5 pt)
noise = torch.randn((hp.batchsize, hp.latent_size), device=device)

# (2) Generate fake images using the generator (hint: you are not supposed to perform gradient
# update on the generator) (1.5 pts)
with torch.no_grad():
    fake_image = generator(noise)

# (3) Calculate the fake loss using the output of the generator (1 pt)
critic_output_fake = critic(fake_image)
critic_loss_fake = critic_output_fake.mean()

# (c) Gradient penalty
alpha = torch.rand((hp.batchsize, 1, 1, 1), device=device)
interpolates = (alpha * real_images + ((1. - alpha)
                * fake_image)).requires_grad_(True)
d_interpolates = critic(interpolates)
gradients = autograd.grad(
    d_interpolates, interpolates, grad_tensor, create_graph=True, only_inputs=True)[0]
gradient_penalty = hp.gp_lambda * (
    (gradients.view(hp.batchsize, -1).norm(dim=1) - 1.) ** 2).mean()

# Implement the aggregated loss using the above three components, be careful with the signs (1 pt)
critic_loss = critic_loss_fake - critic_loss_real + gradient_penalty

critic_loss.backward()
critic_optimizer.step()

if batch_idx % hp.n_critic == 0:
    # Update Generator
    generator_optimizer.zero_grad()

    # Implement the generator loss (2 pts)
    noise = torch.randn((hp.batchsize, hp.latent_size), device=device)
    fake_image = generator(noise)
    critic_output_fake = critic(fake_image)
    generator_loss = -critic_output_fake.mean()

    generator_loss.backward()
    generator_optimizer.step()

# Output training stats
if batch_idx % 100 == 0:
    elapsed_time = time.time() - start_time
    print(f"[{epoch:>2}/{hp.num_epochs}][{iters:>7}][{elapsed_time:8.2f}s]\t"
          f"d_loss/g_loss: {critic_loss.item():4.2}/{generator_loss.item():4.2}\t")

# Save Losses for plotting later
generator_losses.append(generator_loss.item())
critic_losses.append(critic_loss.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 500 == 0) or ((epoch == hp.num_epochs - 1) and (batch_idx == len(dataloader) - 1)):
    with torch.no_grad():
        fake_images = generator(fixed_noise).cpu()
    img_list.append(vutils.make_grid(
        fake_images, padding=2, normalize=True))


iters += 1
```

0/5	0	0.25s	d_loss/g_loss: 2.0/0.17
0/5	100	17.78s	d_loss/g_loss: -5.1/-0.35
0/5	200	35.42s	d_loss/g_loss: -4.0/0.62
0/5	300	52.67s	d_loss/g_loss: -4.9/ 1.2
0/5	400	69.67s	d_loss/g_loss: -6.4/ 1.7
0/5	500	86.62s	d_loss/g_loss: -7.0/ 3.4
0/5	600	103.74s	d_loss/g_loss: -7.2/ 1.8
0/5	700	120.90s	d_loss/g_loss: -6.8/ 2.0
0/5	800	138.00s	d_loss/g_loss: -6.7/ 2.5
0/5	900	155.06s	d_loss/g_loss: -7.1/ 1.5
1/5	937	161.56s	d_loss/g_loss: -7.3/ 2.8
1/5	1037	178.64s	d_loss/g_loss: -7.6/ 3.7
1/5	1137	195.72s	d_loss/g_loss: -7.7/ 3.0
1/5	1237	212.80s	d_loss/g_loss: -7.5/ 2.2
1/5	1337	229.88s	d_loss/g_loss: -8.0/ 2.7
1/5	1437	247.04s	d_loss/g_loss: -7.6/ 2.8
1/5	1537	264.12s	d_loss/g_loss: -7.6/ 1.9
1/5	1637	281.21s	d_loss/g_loss: -7.9/ 3.0
1/5	1737	298.30s	d_loss/g_loss: -7.7/ 2.6
1/5	1837	315.33s	d_loss/g_loss: -1.8/-0.14
2/5	1874	321.83s	d_loss/g_loss: -1.4/0.041
2/5	1974	338.93s	d_loss/g_loss: -1.8/ 1.1
2/5	2074	356.08s	d_loss/g_loss: -1.5/0.78
2/5	2174	373.14s	d_loss/g_loss: -1.5/0.12
2/5	2274	390.19s	d_loss/g_loss: -1.3/-0.24
2/5	2374	407.26s	d_loss/g_loss: -1.4/ 1.7
2/5	2474	424.34s	d_loss/g_loss: -0.98/ 0.6
2/5	2574	441.49s	d_loss/g_loss: -1.1/0.26
2/5	2674	458.55s	d_loss/g_loss: -1.5/0.65
2/5	2774	475.59s	d_loss/g_loss: -1.0/0.59
3/5	2811	482.10s	d_loss/g_loss: -0.042/-0.2
3/5	2911	499.19s	d_loss/g_loss: -1.1/0.62
3/5	3011	516.31s	d_loss/g_loss: -1.6/ 3.3
3/5	3111	533.37s	d_loss/g_loss: -1.2/0.41
3/5	3211	550.49s	d_loss/g_loss: -1.1/-0.27
3/5	3311	567.63s	d_loss/g_loss: -1.6/ 1.0
3/5	3411	584.67s	d_loss/g_loss: -1.4/0.57
3/5	3511	601.81s	d_loss/g_loss: -0.95/-0.13

[3/5][3611][618.87s]	d_loss/g_loss: -0.98/ 1.9
[3/5][3711][635.95s]	d_loss/g_loss: -1.4/0.93
[4/5][3748][642.43s]	d_loss/g_loss: -0.32/-0.75
[4/5][3848][659.49s]	d_loss/g_loss: -1.3/ 1.0
[4/5][3948][676.60s]	d_loss/g_loss: -0.99/-1.1
[4/5][4048][693.70s]	d_loss/g_loss: -1.1/0.24
[4/5][4148][710.76s]	d_loss/g_loss: -0.88/-0.73
[4/5][4248][727.85s]	d_loss/g_loss: -1.3/0.055
[4/5][4348][744.89s]	d_loss/g_loss: -0.73/-0.93
[4/5][4448][762.02s]	d_loss/g_loss: -0.85/0.0035
[4/5][4548][779.11s]	d_loss/g_loss: -1.1/0.76
[4/5][4648][796.18s]	d_loss/g_loss: -0.96/-0.41

Visualization (2 pts)

```
# Visualize the loss
# include the figure in the latex file (1 pt)
plt.title("Generator and Critic Loss During Training")
plt.plot(generator_losses, label="G")
plt.plot(critic_losses, label="D")
plt.xlabel("iterations")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



```
# Visualize the generation (you may scroll to see the animation of training)
# include the final figure in the latex file (1 pt)
import matplotlib.animation as animation
from IPython.display import HTML
# %%capture
fig = plt.figure(figsize=(8, 8))
plt.axis("off")
ims = [[plt.imshow(i.permute(1, 2, 0), animated=True)] for i in img_list]
ani = animation.ArtistAnimation(
    fig, ims, interval=1000, repeat_delay=1000, blit=True)

HTML(ani.to_jshtml())
```

